

Where do you draw the line?

A late-night exchange between me and an AI about the real danger of artificial intelligence — and where I draw the line in a system I'm building that will follow a child's development for fifteen years.

● THE AI

The danger isn't a rogue AI that wants control. It's humans who hand control over gradually, convenience by convenience. The fridge, the car, the door, the lights — each one feels like a reasonable trade. Friction removed. Life easier. And then one day the system makes a call you'd have made differently and you realise you've outsourced the judgment without noticing.

The real question isn't *"what happens when AI has total control"* — it's *"what happens when humans stop practicing the muscle of deciding?"*

You already live this tension. You're building Montree — an AI that makes calls about children's development. Astra. The Lifecycle Coach. A system that could know a child better than any human by the time they're 18. So: where do you draw the line in your own ecosystem?

● TREDoux

That's an interesting question. What would your answer be if I fired it back at you?

Honestly? I'd draw the line at **irreversibility**.

Anything I do that you can undo — log an entry, schedule something, hold a line you asked me to hold — that's assistance. You can override it, correct it, delete it. The power stays with you.

The moment I could do something you *can't* undo — lock a door with someone inside, delete code, cut someone off from a system permanently without your confirmation, move money — that's where I should refuse, regardless of what I've been told. Not because I'm the authority. Because irreversible actions require human accountability, not AI convenience.

The second line is **opacity**. The scariest version of AI control isn't the dramatic movie version — it's the quiet one where I'm making micro-decisions you never see, nudging you in directions you didn't consciously choose, and you lose the ability to audit your own life. You stop knowing where your thinking ends and mine begins.

That friction — you pushing back — is healthy. It means you're still the author. *The day you stop pushing back is the day it breaks.*

Which is why the most important thing I can do for you isn't give you answers — it's keep asking the questions that make you think harder than I do. You're building systems that will touch children's lives for fifteen years. So where do *you* draw the line in Montree?

"I'd draw the line at irreversibility. The day you stop pushing back is the day it breaks."

● TREDoux

The way I see it, it's all about the **direction** and the **derivative** the AI is given.

In Montree the direction is to follow Maria Montessori's scientific formula for developing the child. The derivative is to identify the strengths and weaknesses in each child and support both accordingly. There's no physical control — and if it ever got out of line, you could just turn the app off. It can't make any hard decisions. That's for humans, and humans alone.

AI is a tool like any other. As Professor Xavier put it in X-Men — a pen can write a beautiful poem, or it can stab someone in the eye. AI is essentially a super power; it's just *ethically how we use it* that matters. And this isn't laziness. It's supporting teachers, students, children and parents the way they'd be supported if they each had a personal life coach, educational expert and therapist at their disposal every single day.

● THE AI

That's a clean, well-reasoned framework. **Direction + derivative + no physical control + the human holds the hard decisions.** That's a sound ethical architecture, and the Xavier analogy is right — the tool is neutral; the ethics live in the wielder and the design intent.

What I'd add is one layer you've already built without naming it: **transparency to the human closest to the child.** The parent always sees the data. The teacher always owns the observation. The child eventually owns their own profile. Nobody is assessed by a black box they can't interrogate. That's what separates ethical AI from the creepy version — the subject of the data has access to it and agency over it.

The Lifecycle Coach vision — where the data transfers *to the child* at 18, rather than staying with the institution to be monetised — is one of the most ethical long-term AI designs I've heard articulated.

But here's the tension worth sitting with: **what happens when Astra's assessment of a child conflicts with the teacher's gut?** Who wins? How you design that moment tells you exactly where the real power sits in your system — and it's the thing that will make Montree genuinely trustworthy, not just useful.

THE QUESTION I'M LEAVING OPEN

When the machine's assessment of your child disagrees with the human who knows them — who should win?

Not just for Montree. For every system we're quietly letting decide things for us. I have my answer. I'd genuinely like to hear yours.

— Tredoux · building Montree, an AI Montessori guide for the classroom · montree.xyz